

# Lecture 4

## Sampling:

Inverse Transform, Rejection Sampling, and  
Importance Sampling

## Last Time:

- Expectations and some notation
- The Law of large numbers
- Simulation and Monte Carlo for Integration
- Sampling and the CLT
- Errors in Monte Carlo

# Expectation $E_f[X]$

$$E_f X = \int x dF(x) = \begin{cases} \sum_x x f(x) & \text{if } X \text{ is discrete} \\ \int x f(x) dx & \text{if } X \text{ is continuous} \end{cases}$$

LOTUS, if  $Y = r(X)$ :

$$E[Y] = \int r(x) dF(x)$$

If  $r(X) = I_A(X)$ , Indicator for event  $A$ ,  $p(X \in A) = E_F[I_A(X)] =$   
frequentist probability

# Law of Large numbers (LLN)

- Expectations become sample averages. Convergence for large N.

$$\begin{aligned} E_f[g] &= \int g(x) dF = \int g(x) f(x) dx \\ &= \lim_{n \rightarrow \infty} \frac{1}{N} \sum_{x_i \sim f} g(x_i) \end{aligned}$$

- foundation of Monte Carlo techniques for expectations and integrals, which allow us to replace integration with summation


# Central Limit Theorem

- note that we compute integrals from samples in one replication
- the sample averages are distributed around the true (distribution) expectation in a gaussian distribution with standard error

$$s = \frac{\sigma}{\sqrt{n}}$$

- which mean to use depends on the accuracy you desire

# Monte Carlo $\pi$



- LLN says throw rocks to compute expectation below
- $E_f[I_{\in C}(X, Y)] = \int \int_{\in C} f_{X,Y}(x, y) dx dy$
- which is probability of being in C
- If  $f_{X,Y}(x, y) \sim Uniform(V)$ :

$$= \frac{1}{V} \int \int_{\in C} dx dy = \frac{A}{V}$$

# Formalize Monte Carlo Integration idea

For Uniform pdf:  $U_{ab}(x) = 1/V = 1/(b - a)$

$$J = \int_a^b f(x)U_{ab}(x) dx = \int_a^b f(x) dx / V = I/V$$

From LOTUS and the law of large numbers:


$$I = V \times J = V \times E_U[f] = V \times \lim_{n \rightarrow \infty} \frac{1}{N} \sum_{x_i \sim U} f(x_i)$$

# Today: We need Samples

- to compute expectations, integrals and do statistics, we need samples
- we start that journey today
- inverse transform
- rejection sampling
- importance sampling: a direct, low-variance way to do integrals and expectations



# Inverse transform



# algorithm

The CDF  $F$  must be invertible!

1. get a uniform sample  $u$  from  $Unif(0, 1)$
2. solve for  $x$  yielding a new equation  $x = F^{-1}(u)$  where  $F$  is the CDF of the distribution we desire.
3. repeat.

## Why does it work?

$$F^{-1}(u) = \text{smallest } x \text{ such that } F(x) \geq u$$

What distribution does random variable  $y = F^{-1}(u)$  follow?

The CDF of  $y$  is  $p(y \leq x)$ . Since  $F$  is monotonic:

$$p(y \leq x) = p(F(y) \leq F(x)) = p(u \leq F(x)) = F(x)$$

$F$  is the CDF of  $y$ , thus  $f$  is the pdf.

## Example: exponential

pdf:  $f(x) = \frac{1}{\lambda} e^{-x/\lambda}$  for  $x \geq 0$  and  $f(x) = 0$  otherwise.


$$u = \int_0^x \frac{1}{\lambda} e^{-x'/\lambda} dx' = 1 - e^{-x/\lambda}$$

Solving for  $x$

$$x = -\lambda \ln(1 - u)$$

# code

```
p = lambda x: np.exp(-x)
CDF = lambda x: 1-np.exp(-x)
invCDF = lambda r: -np.log(1-r) # invert the CDF
xmin = 0 # the lower limit of our domain
xmax = 6 # the upper limit of our domain
rmin = CDF(xmin)
rmax = CDF(xmax)
N = 10000
# generate uniform samples in our range then invert the CDF
# to get samples of our target distribution
R = np.random.uniform(rmin, rmax, N)
X = invCDF(R)
hinfo = np.histogram(X,100)
plt.hist(X,bins=100, label=u'Samples');
# plot our (normalized) function
xvals=np.linspace(xmin, xmax, 1000)
plt.plot(xvals, hinfo[0][0]*p(xvals), 'r', label=u'p(x)')
plt.legend()
```



# Box-Muller

- how to draw from a normal?
- the CDF integral is not analytically solvable.

$$I = \frac{1}{2\pi} \int_{-\infty}^x e^{-x'^2/2} dx'$$

- can do numerical inversion (out of scope) or use box-muller trick.
- trick involves starting with two Normals  $N(0, 1)$

$$X \sim N(0, 1), Y \sim N(0, 1) \implies X, Y \sim N(0, 1)N(0, 1)$$

pdf:

$$f_{XY}(x, y) = \frac{1}{\sqrt{2\pi}} e^{-x^2/2} \times \frac{1}{\sqrt{2\pi}} e^{-y^2/2} = \frac{1}{2\pi} \times e^{-r^2/2}$$

where  $r^2 = x^2 + y^2$ .

Using polar co-ordinates  $r$  and  $\theta$ , we have...

$$\Theta \sim \text{Unif}(0, 2\pi), S = R^2 \sim \text{Exp}(1/2)$$

$$s = r^2 = -2\ln(1 - u)$$

$$r = \sqrt{-2\ln(u_1)}, \theta = 2\pi u_2$$

where  $u_1$  and  $u_2 \sim \text{Unif}(0, 1)$ .

Now, use  $x = r \cos\theta, y = r \sin\theta$  to obtain Normal samples.

What is  $f_{R,\Theta}(r, \theta)$ ?



# General transforms of a pdf

Let  $z = g(x)$  so that  $x = g^{-1}(z)$

Define the Jacobian  $J(z)$  of the transformation  $x = g^{-1}(z)$  as the partial derivatives matrix of the transformation.

Then:

$$f_Z(z) = f_X(g^{-1}(z)) \times \det(J(z))$$


Let  $g : r = \sqrt{x^2 + y^2}, \tan(\theta) = y/x$ . Then  $g^{-1} : x = r \cos(\theta),$   
 $y = r \sin(\theta)$

$$J = \begin{pmatrix} \cos(\theta) & \sin(\theta) \\ -r \sin(\theta) & r \cos(\theta) \end{pmatrix}, \det(J) = r$$

$$f_{R,\Theta}(r, \theta) = f_{X,Y}(r \cos(\theta), r \sin(\theta)) \times r$$

$$= \frac{1}{\sqrt{2\pi}} e^{-(r \cos(\theta))^2 / 2} \times \frac{1}{\sqrt{2\pi}} e^{-(r \sin(\theta))^2 / 2} = \frac{1}{2\pi} \times e^{-r^2 / 2} \times r.$$


# Rejection Sampling



- Generate samples from a uniform distribution with support on the rectangle
- See how many fall below  $y(x)$  at a specific  $x$ .

## Algorithm

1. Draw  $x$  uniformly from  $[x_{\min}, x_{\max}]$
2. Draw  $y$  uniformly from  $[0, y_{\max}]$
3. if  $y < f(x)$ , accept the sample
4. otherwise reject it
5. repeat



# example


```
P = lambda x: np.exp(-x)
xmin = 0 # the lower limit of our domain
xmax = 10 # the upper limit of our domain
ymax = 1
#you might have to do an optimization to find this.
N = 10000 # the total of samples we wish to generate
accepted = 0 # the number of accepted samples
samples = np.zeros(N)
count = 0 # the total count of proposals

while (accepted < N):
    # pick a uniform number on [xmin, xmax) (e.g. 0...10)
    x = np.random.uniform(xmin, xmax)
    # pick a uniform number on [0, ymax)
    y = np.random.uniform(0,ymax)
    # Do the accept/reject comparison
    if y < P(x):
        samples[accepted] = x
        accepted += 1

    count +=1

print("Count",count, "Accepted", accepted)
hinfo = np.histogram(samples,30)
plt.hist(samples,bins=30, label=u'Samples');
xvals=np.linspace(xmin, xmax, 1000)
plt.plot(xvals, hinfo[0][0]*P(xvals), 'r', label=u'P(x)')
plt.legend()
```

Count 100294 Accepted 10000



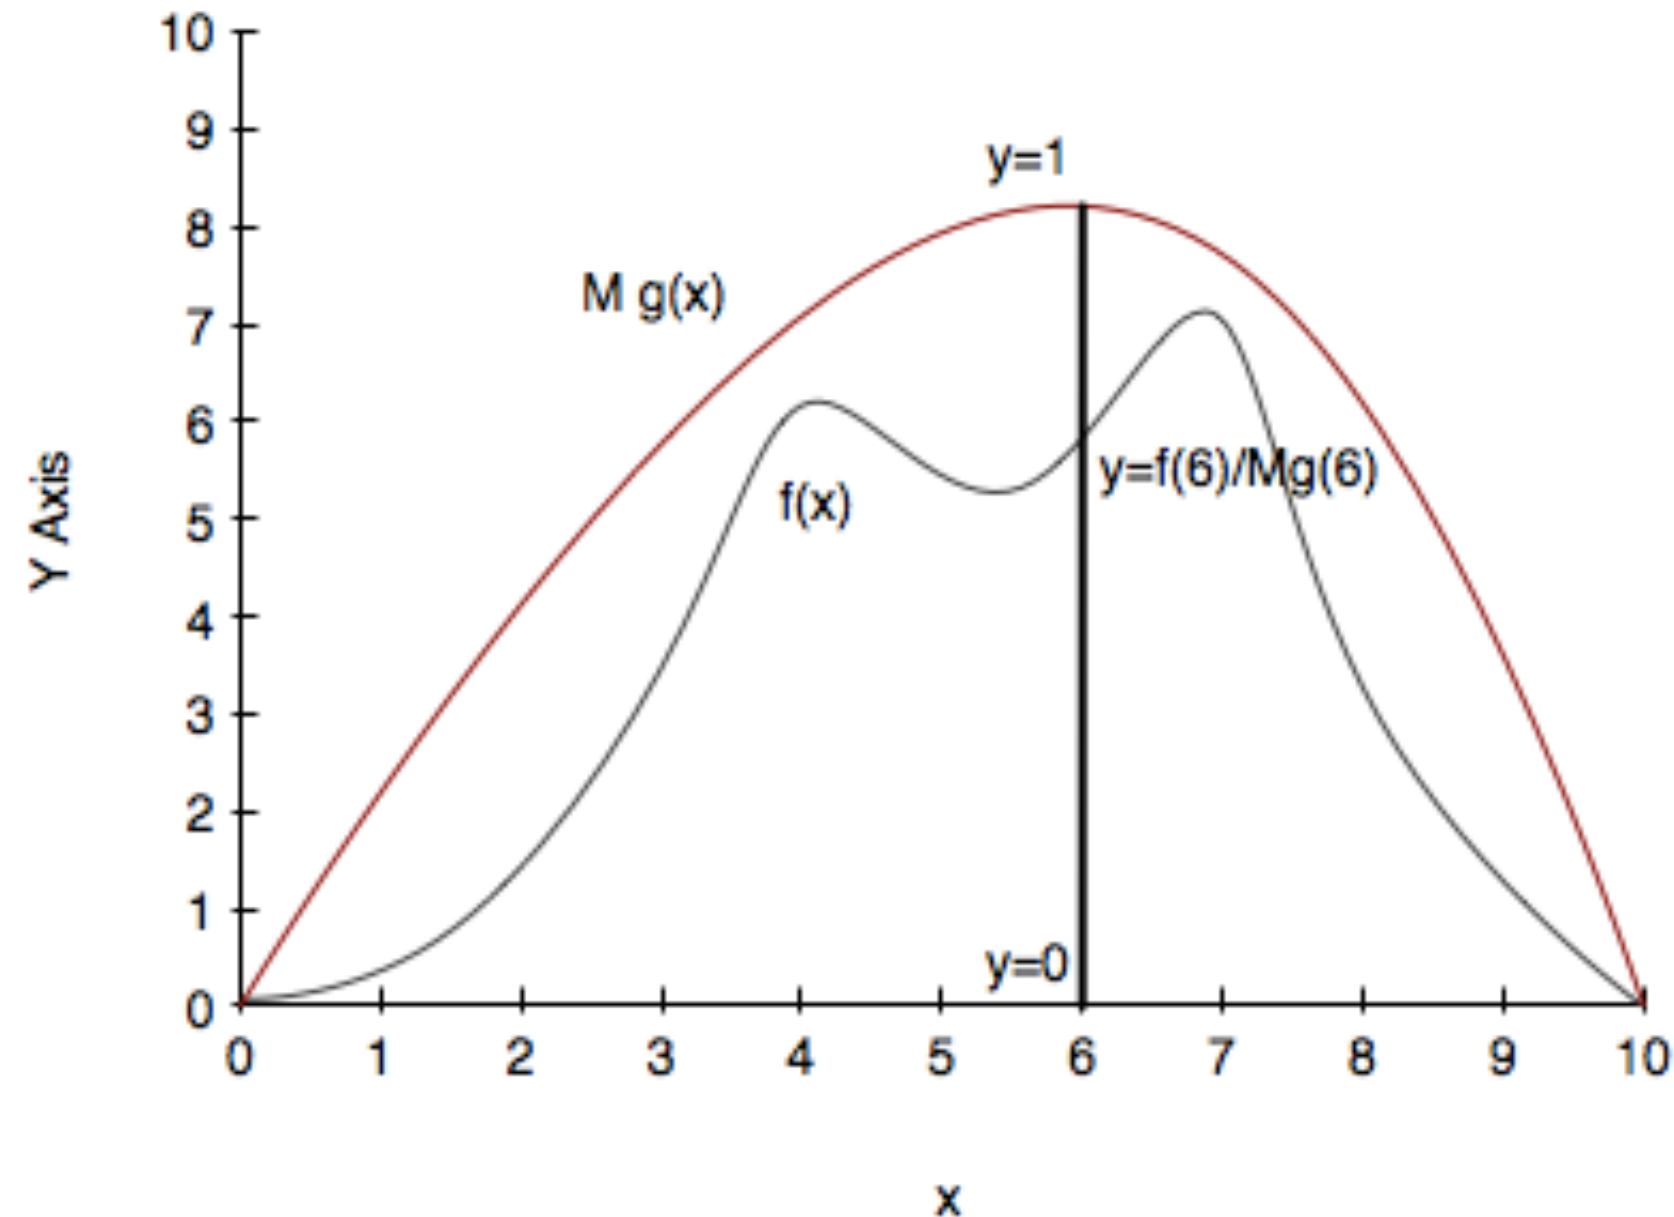
# problems

- determining the supremum may be costly
- the functional form may be complex for comparison
- even if you find a tight bound for the supremum, basic rejection sampling is very inefficient: **low acceptance probability**
- infinite support

# Rejection on steroids


Introduce a **proposal density**  $g(x)$ .

- $g(x)$  is easy to sample from and (calculate the pdf)
- Some  $M$  exists so that  $M g(x) > f(x)$  in your entire domain of interest
- ideally  $g(x)$  will be somewhat close to  $f$
- optimal value for  $M$  is the supremum over your domain of interest of  $f/g$ .
- probability of acceptance is  $1/M$



# Algorithm

1. Draw  $x$  from your proposal distribution  $g(x)$
2. Draw  $y$  uniformly from  $[0,1]$
3. if  $y < f(x)/M g(x)$ , accept the sample
4. otherwise reject it
5. repeat





# Example

```
p = lambda x: np.exp(-x) # our distribution
g = lambda x: 1/(x+1) # our proposal pdf (we're thus choosing M to be 1)
invCDFg = lambda x: np.log(x + 1) # generates our proposal using inverse sampling
xmin = 0 # the lower limit of our domain
xmax = 10 # the upper limit of our domain
# range limits for inverse sampling
umin = invCDFg(xmin)
umax = invCDFg(xmax)
N = 10000 # the total of samples we wish to generate
accepted = 0 # the number of accepted samples
samples = np.zeros(N)
count = 0 # the total count of proposals

while (accepted < N):

    # Sample from g using inverse sampling
    u = np.random.uniform(umin, umax)
    xproposal = np.exp(u) - 1


    # pick a uniform number on [0, 1)
    y = np.random.uniform(0, 1)

    # Do the accept/reject comparison
    if y < p(xproposal)/g(xproposal):
        samples[accepted] = xproposal
        accepted += 1

    count += 1

print("Count", count, "Accepted", accepted)
# get the histogram info
hinfo = np.histogram(samples, 50)
plt.hist(samples, bins=50, label=u'Samples');
xvals=np.linspace(xmin, xmax, 1000)
plt.plot(xvals, hinfo[0][0]*p(xvals), 'r', label=u'p(x)')
plt.plot(xvals, hinfo[0][0]*g(xvals), 'k', label=u'g(x)')
plt.legend()
```

Count 23809 Accepted 10000



# Importance sampling

The basic idea behind importance sampling is that we want to draw more samples where  $h(x)$ , a function whose integral or expectation we desire, is large. In the case we are doing an expectation, it would indeed be even better to draw more samples where  $h(x)f(x)$  is large, where  $f(x)$  is the pdf we are calculating the integral with respect to.

Unlike rejection sampling we use all samples!!

$$E_f[h] = \int_V f(x)h(x)dx.$$


Choosing a proposal distribution  $g(x)$ :

$$E_f[h] = \int h(x)g(x) \frac{f(x)}{g(x)} dV$$

$$E_f[h] = \lim_{N \rightarrow \infty} \frac{1}{N} \sum_{x_i \sim g(\cdot)} h(x_i) \frac{f(x_i)}{g(x_i)}$$

If  $w(x_i) = f(x_i)/g(x_i)$ :

$$E_f[h] = \lim_{N \rightarrow \infty} \frac{1}{N} \sum_{x_i \sim g(\cdot)} w(x_i)h(x_i)$$



# Variance reduction

Usually:  $\hat{V} = \frac{V_f[h(x)]}{N}$

Importance Sampling:  $\hat{V} = \frac{V_g[w(x)h(x)]}{N}$

Minimize  $V_g[w(x)h(x)]$  (make 0), if:

$$w(x)h(x) = C \implies f(x)h(x) = Cg(x), \dots$$

Gives us  $g(x) = \frac{f(x)h(x)}{E_f[h(x)]}$

To get low variance, we must have  $g(x)$  **large where the product  $f(x)h(x)$  is large.**

Or,  $\frac{g(x)}{f(x)}$  ought to be large where  $h(x)$  is large. This means that, as we said earlier, choose more samples near the peak.

# Example: integral of $x \sin(x)$

```
mu = 2;
sig = .7;
f = lambda x: np.sin(x)*x
infun = lambda x: np.sin(x)-x*np.cos(x)
p = lambda x: (1/np.sqrt(2*np.pi*sig**2))*np.exp(-(x-mu)**2/(2.0*sig**2))
normfun = lambda x: norm.cdf(x-mu, scale=sig)
# range of integraion
xmax =np.pi
xmin =0
N =1000 # Number of draws

# Just want to plot the function
x=np.linspace(xmin, xmax, 1000)
plt.plot(x, f(x), 'b', label=u'Original  $x\sin(x)$ ')
plt.plot( x, p(x), 'r', label=u'Importance Sampling Function: Normal')
plt.plot(x, np.ones(1000)/np.pi, 'k')
xis = mu + sig*np.random.randn(N,1);
plt.plot(xis, 1/(np.pi*p(xis)), '.', alpha=0.1)

# IMPORTANCE SAMPLING
Iis = np.zeros(1000)
for k in np.arange(0,1000):
    # DRAW FROM THE GAUSSIAN mean =2 std = sqrt(0.4)
    xis = mu + sig*np.random.randn(N,1);
    xis = xis[ (xis<xmax) & (xis>xmin)] ;
    # normalization for gaussian from 0..pi
    normal = normfun(np.pi)-normfun(0);
    Iis[k] =np.mean(f(xis)/p(xis))*normal;
```

Exact solution is: 3.14159265359

Mean basic MC estimate: 3.14068341144

Standard deviation of our estimates: 0.0617743877206

Mean importance sampling MC estimate: 3.14197268362

Standard deviation of our estimates: 0.0161935244302

