

Lecture 6

Gradient Descent

Last Time:

- Machine learning, especially supervised learning
- Bias, variance, and overfitting
- Minimized an objective function, called error or cost or risk
- Did this on training set, showed test set was a good proxy for out of sample error
- Fit hyper-parameters on validation set

LLN: Expectations \rightarrow sample averages

$$E_f[R] = \int R(x) f(x) dx = \lim_{n \rightarrow \infty} \frac{1}{N} \sum_{x_i \sim f} R(x_i)$$

Empirical Risk Minimization:

$$R_{\mathcal{D}} = E_f[R] \sim \frac{1}{N} \sum_{x_i \sim f} R(x_i)$$


on training set(sample) \mathcal{D} .

Today: optimization using gradient descent

- gradient descent
- stochastic gradient descent

Remember Convex (bowl) like functions have 1 global minimum

Newton's Method



Find a zero of the first derivative.

Gradients and Hessians

$$J(\bar{\theta}) = \theta_1^2 + \theta_2^2$$

$$\text{Gradient: } \nabla_{\bar{\theta}} (J) = \frac{\partial J}{\partial \bar{\theta}} = \begin{pmatrix} 2\theta_1 \\ 2\theta_2 \end{pmatrix}$$

$$\text{Hessian } H = \begin{pmatrix} 2 & 0 \\ 0 & 2 \end{pmatrix}$$

Hessian gives curvature. Why not use it?


Gradient ascent (descent)

basically go opposite the direction of the derivative.


Consider the objective function:

$$J(x) = x^2 - 6x + 5$$


```
gradient = fprime(old_x)
move = gradient * step
current_x = old_x - move
```




good step size



too big step size



too small step size



Example: Linear Regression

$$\hat{y} = f_{\theta}(x) = \theta^T x$$

Cost Function:

$$J(\theta) = \frac{1}{2} \sum_{i=1}^m (f_{\theta}(x^{(i)}) - y^{(i)})^2$$

Gradient Descent

$$\theta := \theta - \eta \nabla_{\theta} J(\theta) = \theta - \eta \sum_{i=1}^m \nabla J_i(\theta)$$


where η is the learning rate.

ENTIRE DATASET NEEDED

```
for i in range(n_epochs):  
    params_grad = evaluate_gradient(loss_function, data, params)  
    params = params - learning_rate * params_grad`
```

Linear Regression: Gradient Descent

$$\theta_j := \theta_j + \alpha \sum_{i=1}^m (y^{(i)} - f_{\theta}(x^{(i)})) x_j^{(i)}$$



Stochastic Gradient Descent

$$\theta := \theta - \alpha \nabla_{\theta} J_i(\theta)$$

ONE POINT AT A TIME

```
for i in range(nb_epochs):  
    np.random.shuffle(data)  
    for example in data:  
        params_grad = evaluate_gradient(loss_function, example, params)  
        params = params - learning_rate * params_grad
```

Mini-Batch: do some at a time

Linear Regression: SGD

$$\theta_j := \theta_j + \alpha(y^{(i)} - f_{\theta}(x^{(i)}))x_j^{(i)}$$