

## Lecture 20

# Model Comparison and glms

# Previously

- glms
- model Checking


# Today

- glms, contd
- oceanic tools example and centering
- model comparison
- oceanic tools and other models model comparison
- poisson over-dispersion and hierarchical modeling

# Back to Poisson GLMs

From Mcelreath:

The island societies of Oceania provide a natural experiment in technological evolution. Different historical island populations possessed tool kits of different size. These kits include fish hooks, axes, boats, hand plows, and many other types of tools. A number of theories predict that larger populations will both develop and sustain more complex tool kits. So the natural variation in population size induced by natural variation in island size in Oceania provides a natural



# Model M1


	culture	population	contact	total_tools	mean_TU	logpop	clevel
0	Malekula	1100	low	13	3.2	7.003065	0
1	Tikopia	1500	low	22	4.7	7.313220	0
2	Santa Cruz	3600	low	24	4.0	8.188689	0
3	Yap	4791	high	43	5.0	8.474494	1
4	Lau Fiji	7400	high	33	5.0	8.909235	1
5	Trobriand	8000	high	19	4.0	8.987197	1
6	Chuuk	9200	high	40	3.8	9.126959	1
7	Manus	13000	low	28	6.6	9.472705	0
8	Tonga	17500	high	55	5.4	9.769956	1
9	Hawaii	275000	low	71	6.6	12.524526	0

$$T_i \sim \text{Poisson}(\lambda_i)$$
$$\log(\lambda_i) = \alpha + \beta_P \log(P_i) + \beta_C C_i + \beta_{PC} C_i \log(P_i)$$
$$\alpha \sim N(0, 100)$$
$$\beta_P \sim N(0, 1)$$
$$\beta_C \sim N(0, 1)$$
$$\beta_{PC} \sim N(0, 1)$$

```
with pm.Model() as m1:
    betap = pm.Normal("betap", 0, 1)
    betac = pm.Normal("betac", 0, 1)
    betapc = pm.Normal("betapc", 0, 1)
    alpha = pm.Normal("alpha", 0, 100)
    loglam = alpha + betap*df.logpop +
              betac*df.clevel + betapc*df.clevel*df.logpop
    y = pm.Poisson("ntools", mu=t.exp(loglam), observed=df.total_tools)

with m1:
    trace=pm.sample(10000, njobs=2)
Average ELBO = -55.784:
100%|██████████| 200000/200000 [00:15<00:00, 13019.16it/s] 12683.03it/s]
100%|██████████| 10000/10000 [01:59<00:00, 83.80it/s]
```

# Posteriors for M1




- traces and autocorrelations look good
- The posterior for  $\beta_p$  tightly constrained, and as expected from theory, shows a positive effect.
- The posteriors for  $\beta_c$  and  $\beta_{pc}$  both overlap 0 substantially, and seem comparatively poorly constrained.
- no substantial effect of contact rate, directly or through the interaction?

# You would be wrong: counterfactual predictions

$\lambda$  traces for high-contact and low contact,  
log(population) of 8.

```
lamlow = lambda logpop: trace['alpha']+trace['betap']*logpop  
lamhigh = lambda logpop: trace['alpha']+(trace['betap'] +  
    trace['betapc'])*logpop + trace['betac']  
sns.distplot(lamhigh(8) - lamlow(8));
```

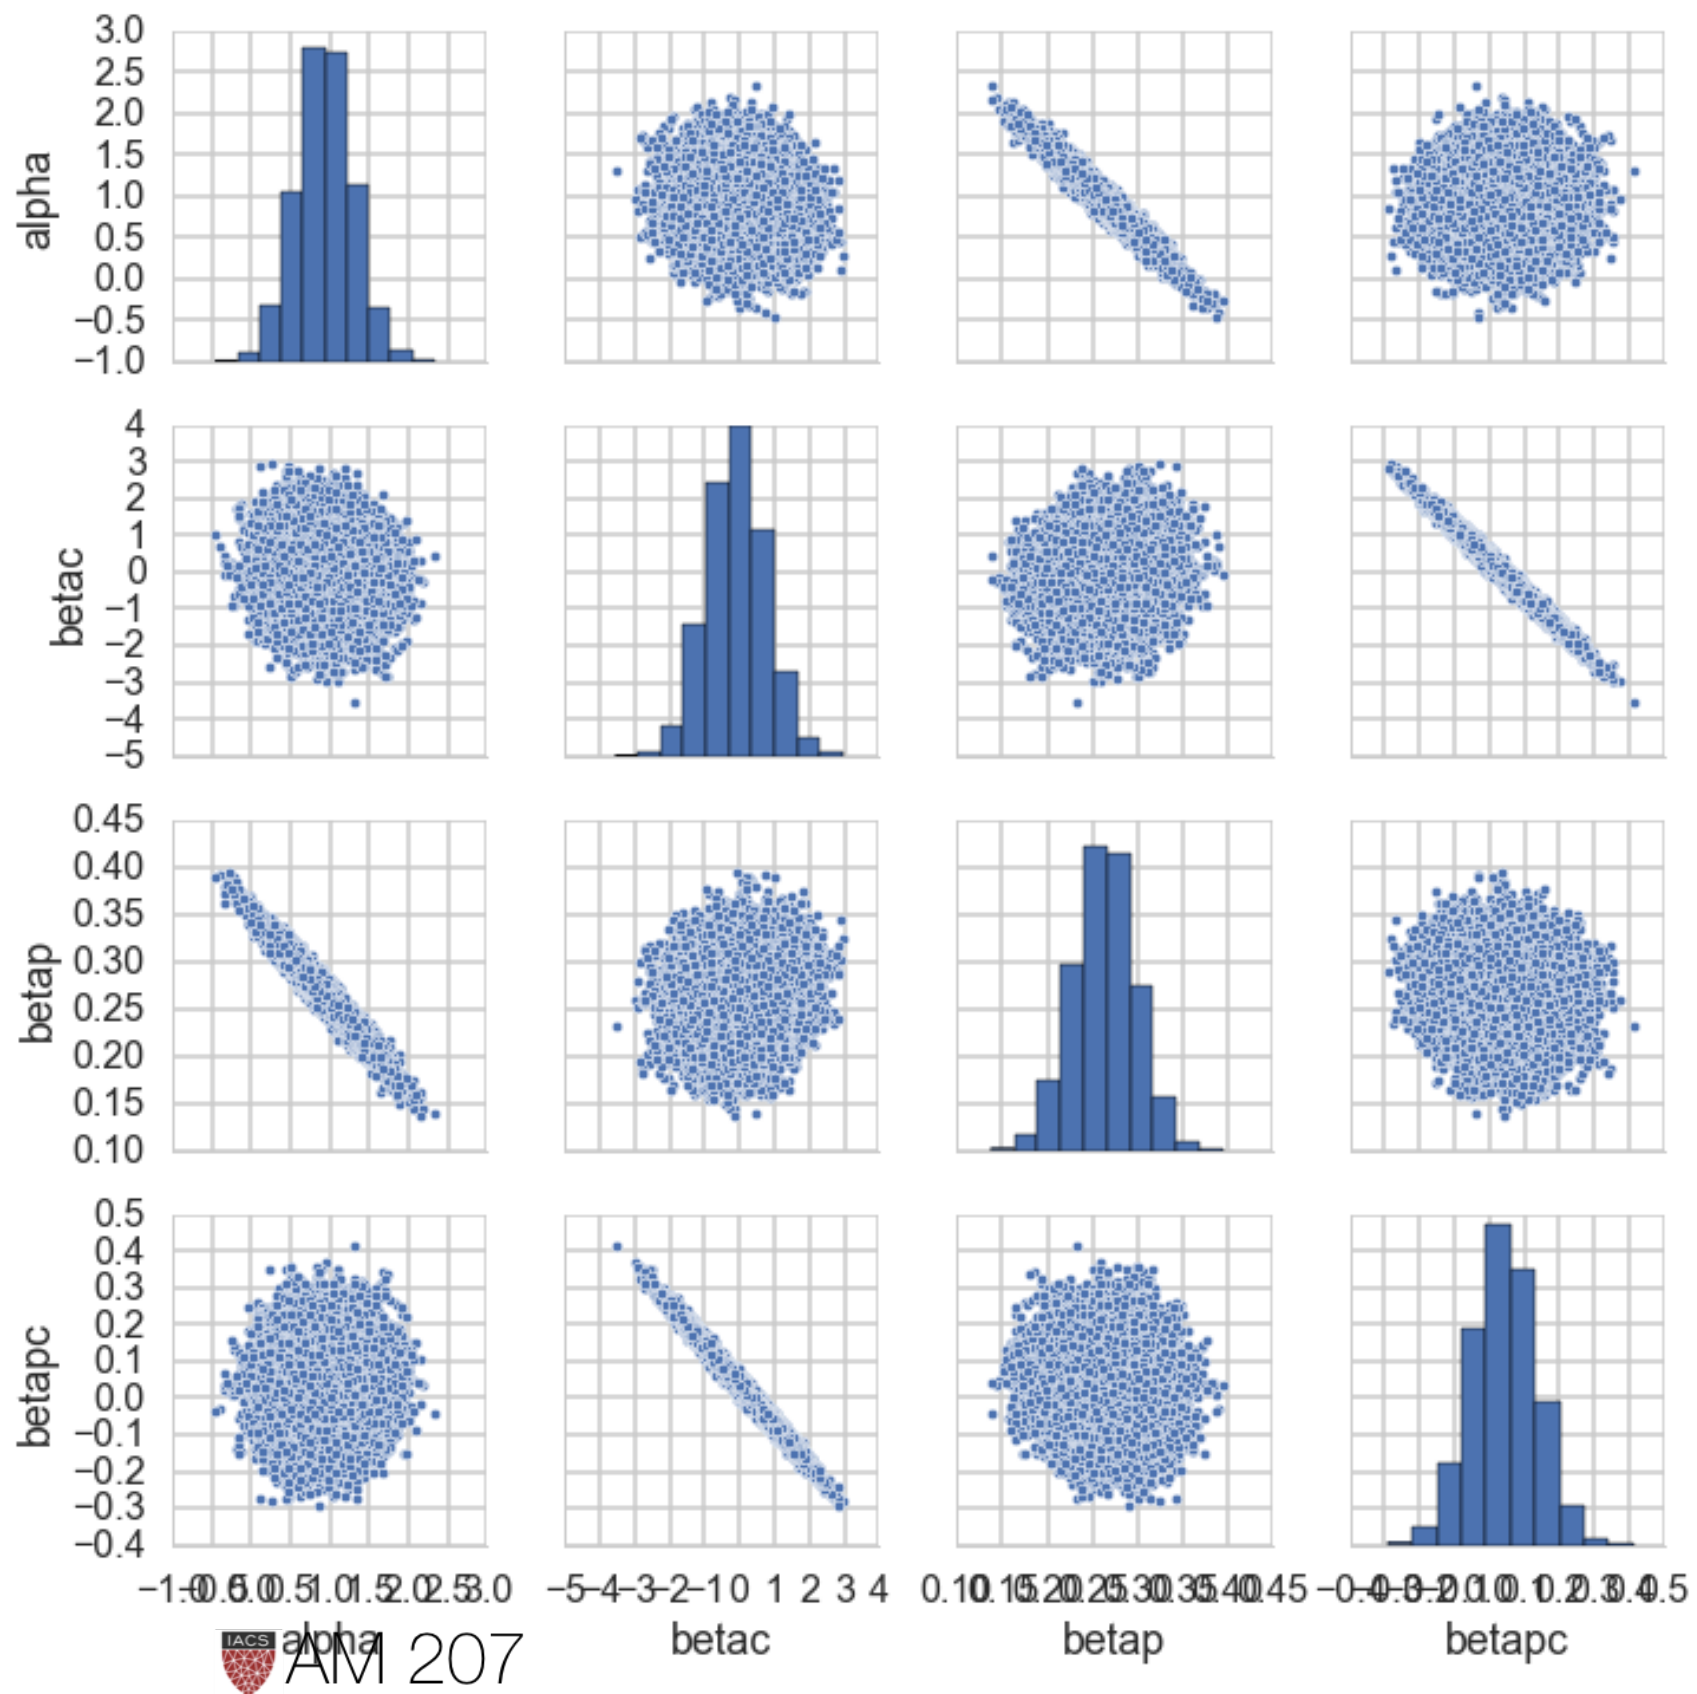
A new kind of model checking.





# What happened?

- very strong negative correlations between  $\alpha$  and  $\beta_p$
- very strong negative correlations between  $\beta_c$  and  $\beta_{pc}$ .
- The latter is the cause for the 0-overlaps.
- When  $\beta_c$  is high,  $\beta_{pc}$  must be low, and vice-versa. Look at the joint uncertainty of the correlated variables rather than just marginals





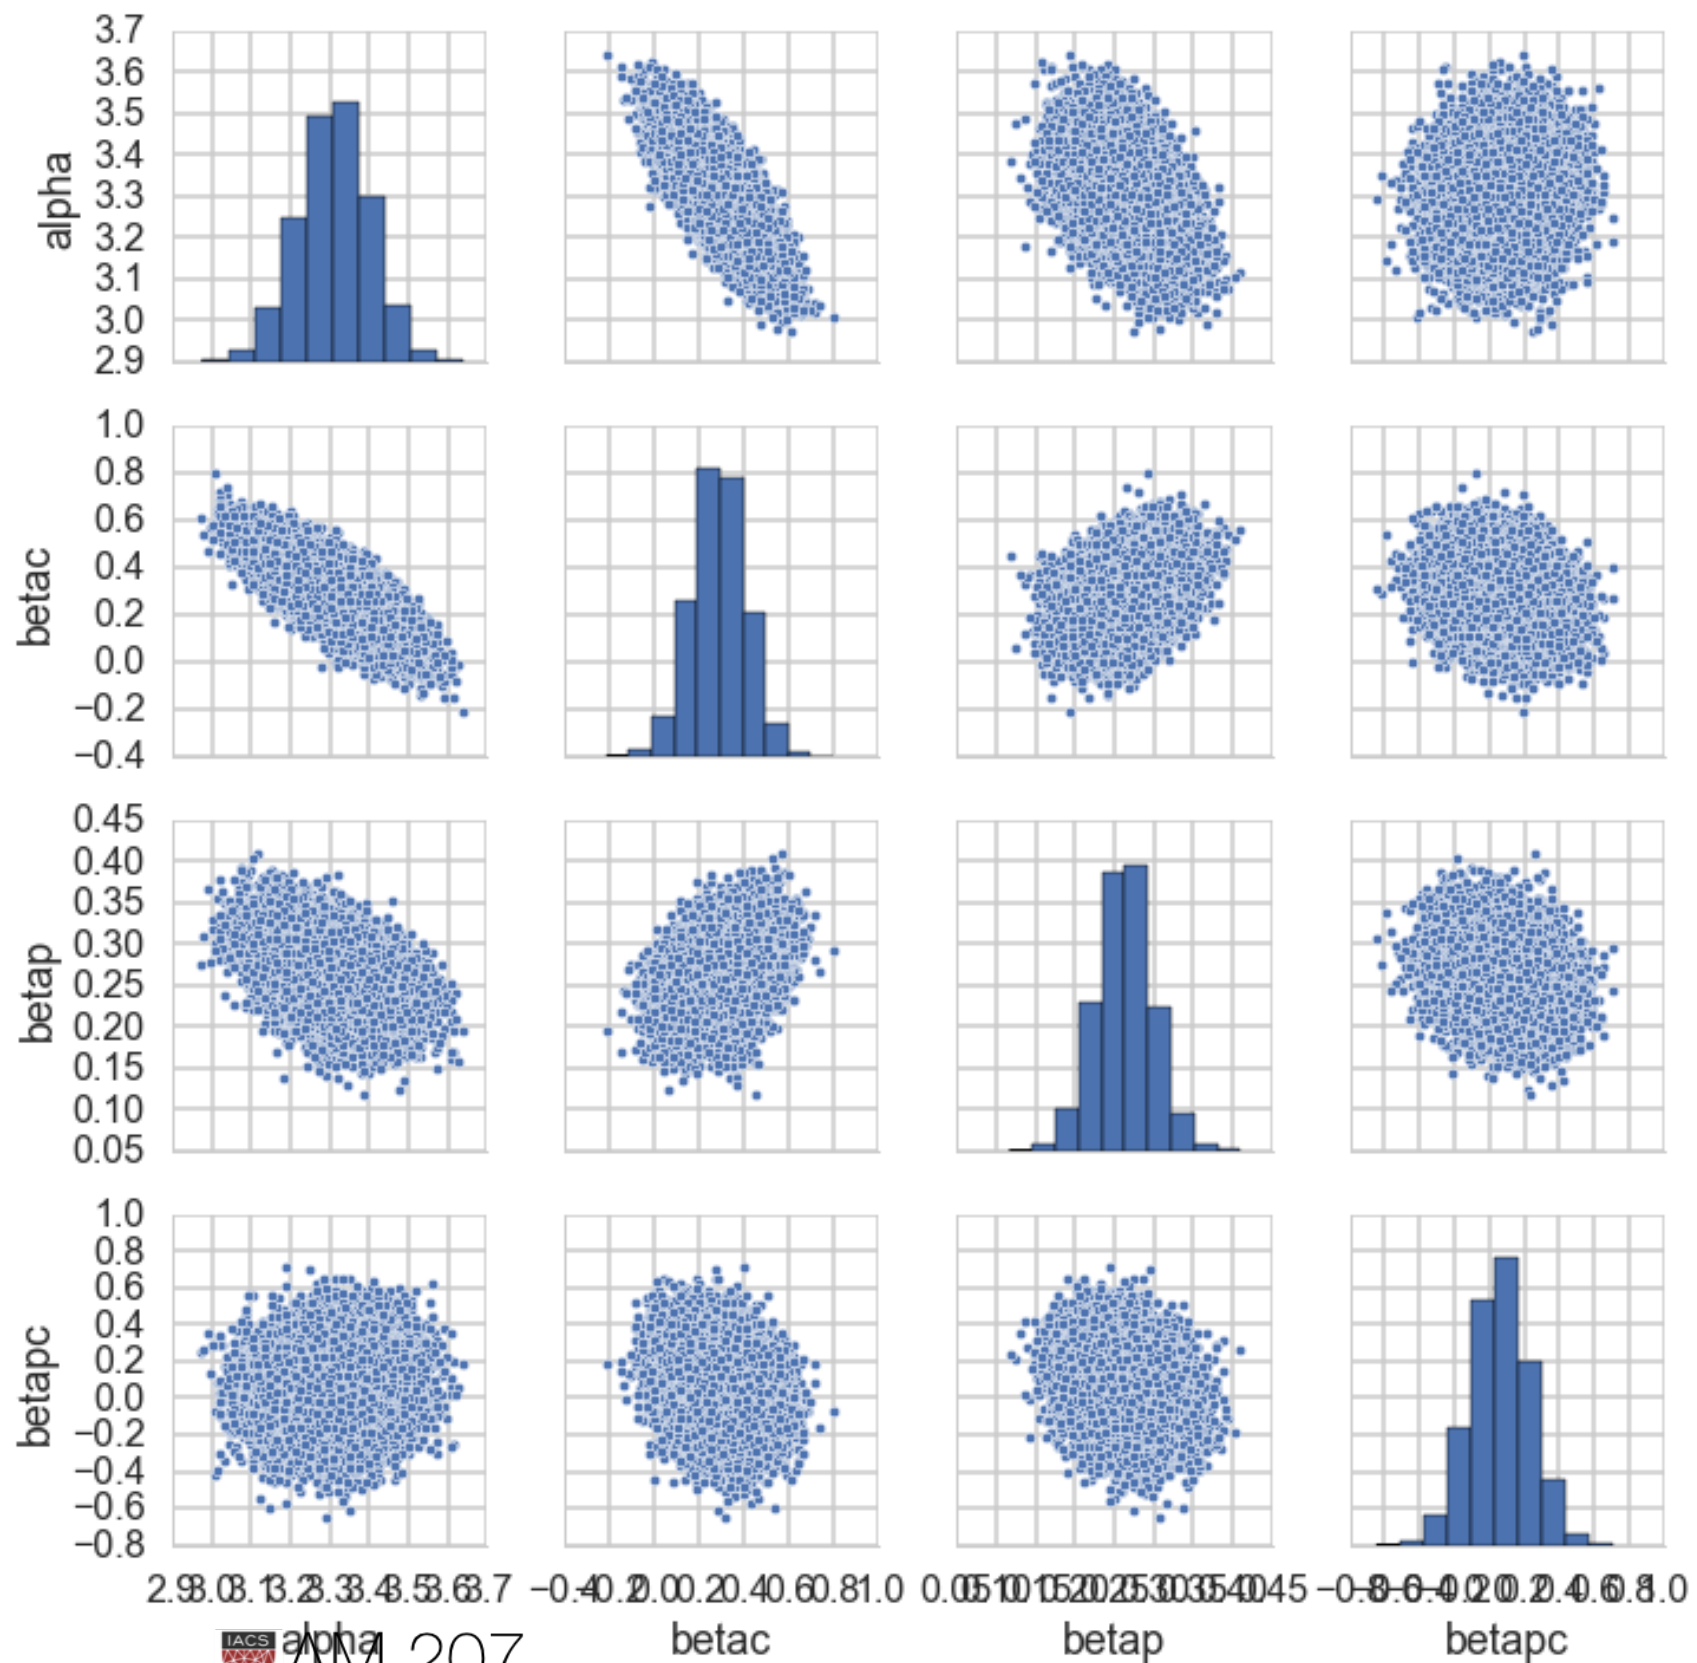
# Fix by centering

- you would have seen the problem in  $n_{eff}$ :


```
{ 'alpha': 8110.0, 'betac': 4600.0, 'betap': 8016.0, 'betapc': 4597.0 }
```

```
with pm.Model() as m1c:
    betap = pm.Normal("betap", 0, 1)
    betac = pm.Normal("betac", 0, 1)
    betapc = pm.Normal("betapc", 0, 1)
    alpha = pm.Normal("alpha", 0, 100)
    loglam = alpha + betap*df.logpop_c + betac*df.clevel + betapc*df.clevel*df.logpop_c
    y = pm.Poisson("ntools", mu=t.exp(loglam), observed=df.total_tools)
```

```
{ 'alpha': 7978.0, 'betac': 7898.0, 'betap': 13621.0, 'betapc': 17703.0 }
```




- better constrained, less correlated, sampling faster and better
- clear effect of contact, effect of interaction not clear yet
- will use model comparison next time for this!



# Model comparison

# Dataset



- 20 data points
- keep train and test, we'll see these later
- center  $x$  for speed of sampling
- fit multiple order polynomials
- The `logp` method of a node gives us the conditional (log) probability of the node given its parents.
- Conditional probability of the observed node, which is also the model node, ...

# model.logp({paramdict})

**gives us likelihood of data  $p(y|etc)$**

```
ridge=3
sigma=0.2
pf=PolynomialFeatures(msize, include_bias=False).fit_transform(xtrain.reshape(-1,1))
print(pf.shape)
with pm.Model() as m:
    alpha = pm.Normal('alpha', 0, 100)
    beta = pm.Normal('beta', mu=0, sd=ridge, shape=msize)
    mu = alpha + pm.math.dot(pf, beta)
    o = pm.Normal('o', mu, sigma, observed=ytrain)
    trace=pm.sample(5000, init='MAP')

meanpoint={'alpha': trace['alpha'].mean(), 'beta': trace['beta'].mean(0)}
{'alpha': 13.840500632702547, 'beta': array([-0.99418043,  2.31483997])}
m.logp(meanpoint), mtest.logp(meanpoint)
(array(-1.699085634529018), array(-5.251454815218393))
```

# Information criteria

- simulate an ensemble of polynomials
- use information criteria to decide between them
- these come from the deviance

$$D_{KL}(p, q) = E_p[\log(p) - \log(q)] = E_p[\log(p/q)] = \sum_i p_i \log\left(\frac{p_i}{q_i}\right) \text{ or } \int dP \log\left(\frac{p}{q}\right)$$

Use **law of large numbers** to replace the true distribution by

its empirical estimate, then we have:

$$D_{KL}(p, q) = E_p[\log(p/q)] = \frac{1}{N} \sum_i (\log(p_i) - \log(q_i))$$

Thus minimizing the KL-divergence involves maximizing  $\sum_i \log(q_i)$ ,  
justifies the maximum likelihood principle.

$$D_{KL}(p, q) - D_{KL}(p, r) = E_p[\log(r) - \log(q)] = E_p[\log(\frac{r}{q})]$$



# Deviance

$$D(q) = -2 \sum_i \log(q_i),$$

then

$$D_{KL}(p, q) - D_{KL}(p, r) = \frac{2}{N} (D(q) - D(r))$$

More generally:  $D(q) = -\frac{N}{2} E_p[\log(q)]$

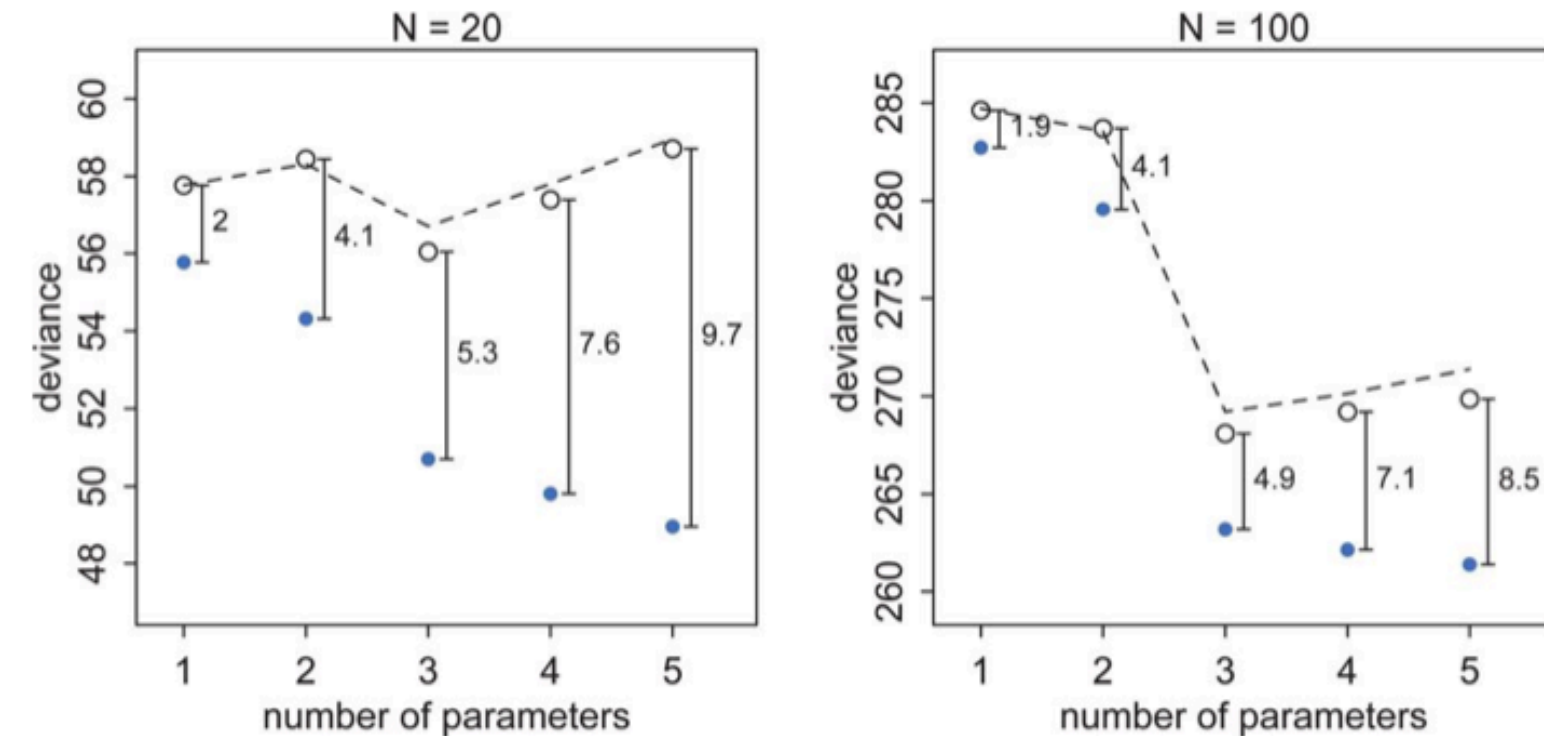
# AIC

Akaike Information Criterion, or AIC:

$$AIC = D_{train} + 2p$$

$$D_{train} = -2 * \log(p(y|\theta_{mle}))$$

- multivariate gaussian posterior
- flat priors
- data >> parameters



# DIC

- uses the posterior distribution, calculable from MCMC.
- multivariate gaussian posterior distribution.

$$D_{train} = -2 * \log(p(y|\theta_{postmean}))$$

$$DIC = D_{train} + 2p_D \text{ where}$$

$$p_{DIC} = 2 * (\log(p(y|\theta_{postmean})) - E_{post}[\log(p(y|\theta))]) \text{ (by monte carlo)}$$

alternative fomulation for  $p_D$ , guaranteed to be positive, is

$$p_D = 2 * Var_{post} [\log(p(y|\theta_{postmean}))]$$

For Model with 19 slope, dic is 114.8922879773209

For Model with 1 slope, dic is 14.754881277586044

For Model with 2 slope, dic is 9.315532339708351

For Model with 3 slope, dic is 14.429626596917327

For Model with 4 slope, dic is 23.634453163719378

For Model with 5 slope, dic is 26.99865973609573

For Model with 10 slope, dic is 60.436331238741445

For Model with 15 slope, dic is 92.68228005183747

# Bayesian deviance

$D(q) = -\frac{N}{2} E_p[\log(pp(y))]$  posterior predictive for points  $y$  on the test set or future data

replace joint posterior predictive over new points  $y$  by product of marginals:

$$\text{ELPD: } \sum_i E_p[\log(pp(y_i))]$$

Since we do not know the true distribution  $p$ ,

replace elpd:  $\sum_i E_p[\log(p(y_i))]$

by the computed "log pointwise predictive density" (lppd) **in-sample**

$$\sum_j \log \langle p(y_j | \theta) \rangle = \sum_j \log \left( \frac{1}{S} \sum_s p(y_j | \theta_s) \right)$$

# WAIC

$$WAIC = lppd + 2p_W$$

where

$$p_W = 2 \sum_i (\log(E_{post}[p(y_i | \theta)]) - E_{post}[\log(p(y_i | \theta))])$$

Once again this can be estimated by

$$\sum_i Var_{post}[\log(p(y_i | \theta))]$$



# Oceanic tools

## Use WAIC to compare models

m2c\_onlyic:  $\text{loglam} = \alpha$

m2c\_onlyc:  $\text{loglam} = \alpha + \text{betac} * \text{df.clevel}$

m2c\_onlyp:  $\text{loglam} = \alpha + \text{betap} * \text{df.logpop}_c$

m2c\_nopc:  $\text{loglam} = \alpha + \text{betap} * \text{df.logpop}_c + \text{betac} * \text{df.clevel}$

m1c:  $\text{loglam} = \alpha + \text{betap} * \text{df.logpop}_c + \text{betac} * \text{df.clevel} + \text{betapc} * \text{df.clevel} * \text{df.logpop}_c$


# Centered

	WAIC	pWAIC	dWAIC	weight	SE	dSE	warning
name							
m2c_nopc	79.3591	4.39013	0	0.846327	11.0543	0	1
m1c	83.8617	6.94776	4.50259	0.0890866	12.2027	4.00079	1
m2c_onlyp	84.5049	3.77558	5.14581	0.0645862	8.91335	19.3282	1
m2c_onlyic	141.327	8.10745	61.9681	2.96038e-14	31.6664	339.158	1
m2c_onlyc	152.975	18.1559	73.6157	8.75158e-17	46.6488	679.109	1

- dWAIC is the difference between each WAIC and the lowest WAIC.
- SE is the standard error of the WAIC estimate.
- dSE is the standard error of the difference in WAIC between each model and the top-ranked model.

$$w_i = \frac{\exp(-\frac{1}{2}dWAIC_i)}{\sum_j \exp(-\frac{1}{2}dWAIC_j)}$$

read each weight as an estimated probability that each model will perform best on future data.




# Uncentered

	WAIC	pWAIC	dWAIC	weight	SE	dSE	warning
name							
m2_nopc	79.1059	4.22647	0	0.61959	11.0612	0	1
m1	80.3046	5.03686	1.19871	0.340258	11.3985	0.571957	1
m2_onlyp	84.5787	3.84888	5.47276	0.0401523	8.98146	20.1717	1
m2_onlyic	141.327	8.10745	62.2212	1.90956e-14	31.6664	338.568	1
m2_onlyc	152.975	18.1559	73.8689	5.64512e-17	46.6488	678.014	1


interaction is overfit. centering decorrelates

# Counterfactual Posterior predictive




# Ensembling

- use WAIC based akaike weights for top 3
- regularizes down the green band at high population by giving more weight to the no-interaction model.



# Overdispersion for only p


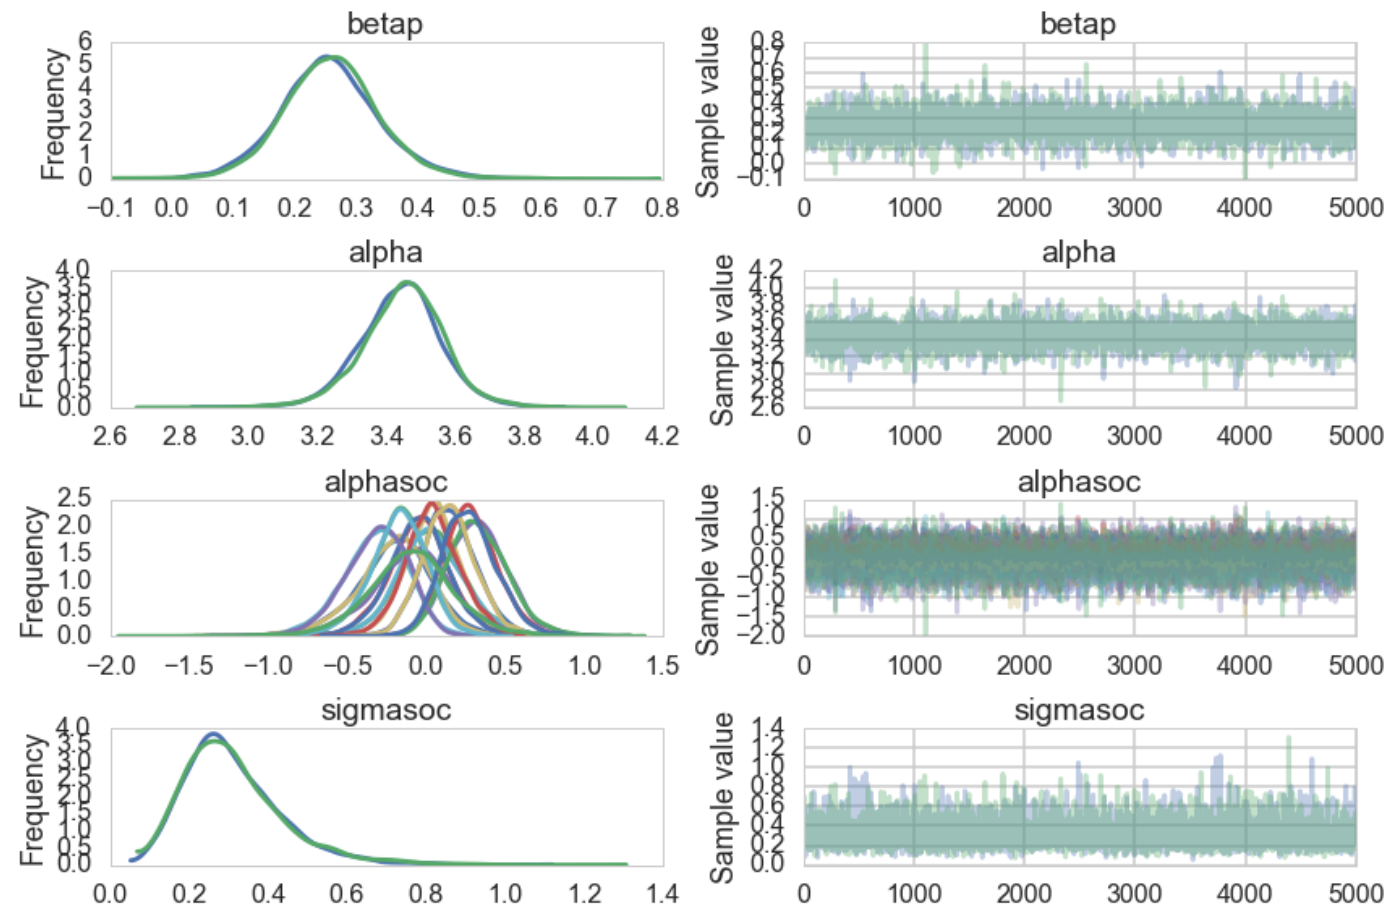




# Varying hierarchical intercepts model

```
with pm.Model() as m3c:
    betap = pm.Normal("betap", 0, 1)
    alpha = pm.Normal("alpha", 0, 100)
    sigmasoc = pm.HalfCauchy("sigmasoc", 1)
    alphasoc = pm.Normal("alphasoc", 0, sigmasoc, shape=df.shape[0])
    loglam = alpha + alphasoc + betap*df.logpop_c
    y = pm.Poisson("ntools", mu=t.exp(loglam), observed=df.total_tools)
```

# Hierarchical Model Posterior predictive



much wider, includes data areas